



NOVA SCHOOL OF
SCIENCE & TECHNOLOGY

NOVA UNIVERSITY OF LISBON

MSC IN COMPUTER SCIENCE

PERFORMANCE COMPARISON BETWEEN DBMSs UNDER TPROC-C WORKLOADS

José Costa (62637)
Rodrigo Albuquerque (70294)
Rodrigo Silva (70567)

DATABASES SYSTEMS

JUNE 9, 2025

Contents

| | | |
|----------|---------------------------------------|-----------|
| 1 | Introduction | 1 |
| 2 | Overview of HammerDB | 1 |
| 2.1 | Overview of TPROC-C | 1 |
| 2.2 | TPROC-C vs TPROC-H | 1 |
| 3 | Problem & DBMS Summary | 1 |
| 4 | Benchmark Description | 2 |
| 4.1 | Benchmark Goals | 2 |
| 4.2 | Test Parameters | 2 |
| 4.3 | Execution Environment | 2 |
| 4.4 | Metrics Collected | 3 |
| 4.5 | Limitations | 3 |
| 5 | Methodology | 3 |
| 5.1 | Hardware and Software Setup | 3 |
| 5.2 | Database Setup | 4 |
| 6 | Results | 4 |
| 7 | Discussion | 4 |
| 8 | Conclusions | 4 |
| 9 | Bibliography | 5 |
| A | TCL Scripts | 6 |
| A.1 | MySQL TCL Script | 6 |
| A.2 | MariaDB TCL Script | 7 |
| A.3 | PostgreSQL TCL Script | 8 |
| B | Docker | 9 |
| B.1 | Docker Compose File | 9 |
| B.2 | Dockerfile for MySQL | 10 |
| B.3 | Dockerfile for MariaDB | 10 |
| C | Configs | 11 |
| C.1 | MySQL Configuration File | 11 |

| | |
|--|-----------|
| C.2 MariaDB Configuration File | 12 |
| C.3 PostgreSQL Configuration File | 14 |
| D Runner Scripts | 15 |
| D.1 Batch Script | 15 |
| D.1.1 PCStats Script | 16 |
| D.2 Shell Script | 17 |
| D.2.1 PCStats Script | 18 |

1 Introduction

This project aims to benchmark and compare the performance of different database systems using the TPC-C workload, a standard for evaluating OLTP (Online Transaction Processing) environments.

Automated scripts are used to run identical tests across multiple databases, each configured with similar settings to ensure a fair comparison.

Key metrics such as Transactions Per Minute (TPM) are measured under varying levels of concurrent users.

The results help identify the strengths and limitations of each database in handling transactional workloads.

In this project, the databases used were PostgreSQL, MySQL and MariaDB.

In total, we ran 54 tests:

- 4 tests scaling the number of virtual users (2 4 8 12) and warehouses (VU*5) on all PCs and databases (48 tests in total);
- 1 test with the number of virtual users set to the same as the number of threads in that PC, warehouses set to VU*5, *allwarehouse* = true on all databases on just one PC (3 test in total);
- 1 test with the number of virtual users set to the same as the number of threads in that PC, warehouses set to VU*5, with default config on all databases on just one PC (3 test in total).

2 Overview of HammerDB

HammerDB is a free, open-source tool for benchmarking the performance of relational databases [1].

It supports popular databases like Oracle, SQL Server, PostgreSQL, MySQL, and more. HammerDB uses industry-standard workloads such as TPROC-C and TPROC-H to simulate real-world database activity.

It offers both a graphical interface and command-line options, making it suitable for developers, DBAs, and system administrators to test, compare, and tune database performance.

In some cases we used HammerDB in docker containers to run the tests, which allows for easy setup and isolation of the testing environment.

In another case, we used the Windows version of HammerDB to run the tests on a Windows machine.

2.1 Overview of TPROC-C

TPROC-C is a benchmark designed to evaluate the performance of database management systems (DBMS) using a transactional workload. It simulates a typical online transaction processing (OLTP) environment, focusing on operations like inserts, updates, and deletes across multiple tables.

2.2 TPROC-C vs TPROC-H

TPROC-H is a benchmark designed for data warehousing and analytical workloads, while TPROC-C is focused on transactional processing. TPROC-H emphasizes complex queries and large data sets, whereas TPROC-C simulates real-time transactions with a focus on insert, update, and delete operations.

3 Problem & DBMS Summary

Modern applications progressively depend on robust, scalable database systems to effectively manage workloads in a transactional manner. Choosing the right Database Management System (DBMS) is critical to achieve the best possible performance, especially where concurrency is high and the workload is varied. With so many DBMSs to select from, with their own strengths, configurations, and community support, getting the right one can be problematic.

The main objective of this study is to provide a clear comparison of how three of the most popular open-source DBMSs like PostgreSQL, MySQL, and MariaDB perform under TPROC-C workloads. This

is a close approximation of OLTP environments and thus is suitable to use in evaluating systems for high-throughput transaction processing.

A brief overview of the DBMSs under test is provided below:

- **PostgreSQL:** It is renowned for support of sophisticated query capabilities, extensibility, and strict support for data integrity features.
- **MySQL:** Extensively used in web development, valued as being easy to use and swift, with excellent ecosystem and support.
- **MariaDB:** A fork of MySQL with an emphasis on improved speed, open development, and additional storage engines.

Each DBMS was installed with identical hardware and software configurations to provide a level playing field for the tests. Experiments were designed to highlight differences in how each system handles transaction loads, concurrency, and configuration parameters. Through comparison of performance in a systematic manner across controlled tests, this study aims to guide database selection on the grounds of empirical evidence and not assumptions.

4 Benchmark Description

The benchmark used in the research is the TPROC-C workload using HammerDB. TPROC-C is designed to simulate an OLTP environment in the average case and is composed of transactions containing new orders, orders to make payment, checking orders status, delivering orders, and updating stock status. The aforementioned operations can be likened to real-world application environments.

4.1 Benchmark Goals

The main objectives of the benchmark are to:

- Compare the throughputs of the DBMSs at different levels of concurrency in terms of Transactions Per Minute, or TPM.
- Monitor the scalability of the systems as virtual users and additional warehouses get created.
- Measure consistency across repeated experiments as well as with different setups.

4.2 Test Parameters

To ensure consistency and fairness, the same configuration template was used for all tests, with the sole variations being the number of virtual users, warehouses, and the specific DBMS being tested. The significant parameters were:

- **Virtual Users (VU):** Simulated clients making postings simultaneously. Applying the values 2, 4, 8, and 12 to perform the scaling test.
- **Warehouses:** A TPROC-C scale unit. Five times the number of virtual users.
- **Test Duration:** Each test was executed in medium one hour.
- **Ramp-Up and Cool-Down Time:** Confirmed with HammerDB setup to allow systems to reach steady state before measurement.

4.3 Execution Environment

All the testing was automated with custom scripts offering the same setup routines and execution across the systems. The environment included combinations of Windows and Linux systems, based on the setup. HammerDB was run in some cases in Docker containers to offer isolation to the test environment and ensure repeatability.

4.4 Metrics Collected

A critical measurement during the benchmarking was the Transaction Per Minute (TPM), as posted by HammerDB. The measurement is indicative of the system capability to process new orders and follow-up transactions in one minute. Secondary measures included CPU usage and memory use.

4.5 Limitations

Even though control over variables, consistency across runs, and bias to a variable as limited as possible was the focus of this testing, there are several limitations:

- Variability in network latency, and relative performance of hardware and storage may be negligible in each of the testing environments.
- Unless otherwise stated, default DBMS tuning parameters were used that are not reflective of, nor necessarily represent, the best performance that could be achieved for each database management system.
- The intent and focus were on relative performance under a specified workload, not full optimization for each system.

5 Methodology

5.1 Hardware and Software Setup

| PC | 1 | 2 | 3 | 4 |
|-----------|------------------|-----------------|-----------------|---------------|
| OS | Windows 11 | Windows 11 | Linux (Unraid) | MacOS Sequoia |
| CPU | AMD Ryzen 5 3600 | Intel i7-13700H | Intel i3-10100F | Apple M1 |
| Cores | 6 | 14 (6P 8E) | 4 | 8 |
| Threads | 12 | 20 | 8 | 8 |
| RAM | 16GB | 16GB | 32GB | 16GB |
| Disk | SSD M.2 NVMe | SSD M.2 NVMe | SSD M.2 NVMe | SSD M.2 NVMe |
| Read | 2500 MB/s | 3500 MB/s | 3500 MB/s | 3400 MB/s |
| Write | 2100 MB/s | 2700 MB/s | 3300 MB/s | 2800 MB/s |
| Test type | Bare metal | Docker | Docker | Docker |

Table 1: Hardware used in the benchmarks

The benchmark tests were run on four different PCs that each met the same specifications for the purpose of consistent performance measurement.

- **PC 1:** All components, including HammerDB and the database systems, were installed directly on the host operating system (bare metal). This installation aimed to assess performance without the overhead associated with running in a container.
- **PCs 2, 3, and 4:** In this case, HammerDB and every DBMS were run within Docker containers. This offered a consistent method that provided isolation between services while ensuring consistency across machines.

We used as our orchestration tool [Docker Compose](#), allowing us to apply and deploy the same setup as configured in multiple machines with minimal processing.

For [MySQL](#) and [MariaDB](#), we created custom Dockerfiles to support the loading of external configuration files. This is because the databases enforced read-only access to the configuration file by default, unlike PostgreSQL, inside the container.

5.2 Database Setup

6 Results

7 Discussion

8 Conclusions

9 Bibliography

- [1] Wikipedia contributors. *HammerDB — Wikipedia, The Free Encyclopedia*. [Online; accessed 25-May-2025]. 2025. URL: <https://en.wikipedia.org/w/index.php?title=HammerDB&oldid=1275860580>.

A TCL Scripts

A.1 MySQL TCL Script

Listing 1: TCL script for MySQL

```
1  #!/usr/bin/tclsh
2  # Set Database & Benchmark
3  dbset db mysql
4  dbset bm TPC-C
5
6  # DB configs
7  diset connection mysql_host 172.22.0.3
8  diset connection mysql_port 3306
9  diset connection mysql_socket "/var/run/mysqld/mysqld.sock"
10 diset tpcc mysql_user root
11 diset tpcc mysql_pass 1234
12 diset tpcc mysql_dbase tpcc
13
14 # Default for WH and VU
15 diset tpcc mysql_count_ware 50
16 diset tpcc mysql_num_vu 10
17
18 # Driver script options
19 diset tpcc mysql_timeprofile true
20 diset tpcc mysql_async_scale false
21 diset tpcc mysql_driver timed
22 # Ensure test is limited by time
23 diset tpcc mysql_total_iterations 1000000
24 # Timed duration
25 diset tpcc mysql_rampup 2
26 diset tpcc mysql_duration 8
27 # Distribute load
28 diset tpcc mysql_allwarehouse false
29
30 # Transactions options
31 tcset refreshrate 10
32 tcset logtotemp 1
33 tcset unique 1
34 tcset timestamps 1
35 # Run
36 foreach z {2 4 8 12} {
37     set w [expr {$z * 5}]
38     diset tpcc mysql_count_ware $w
39     puts "Building Schema for $z TEST"
40     # Delete possible previous data
41     deleteschema
42     vudestroy
43     # Build and load
44     buildschema
45     vudestroy
46     loadscript
47     # Vuser options
48     vuset delay 500
49     vuset repeat 500
50     vuset iterations 1
51     vuset showoutput 0
52     vuset logtotemp 1
53     vuset unique 1
54     vuset nobuff 0
55     vuset timestamps 1
56     puts "Starting $z VU TEST"
57     tcstart
58     vuset vu $z
59     vucreate
60     vurun
61     puts "Waiting 1 for cleanup and collection"
62     after 60000
63     puts "Destroying VU"
64     vudestroy
65     tcstop
66 }
```

A.2 MariaDB TCL Script

Listing 2: TCL script for MariaDB

```
1  #!/usr/bin/tclsh
2  # Set Database & Benchmark
3  dbset db maria
4  dbset bm TPC-C
5
6  # DB configs
7  diset connection maria_host 172.22.0.4
8  diset connection maria_port 3306
9  diset connection mysql_socket "/var/run/mysqld/mysqld.sock"
10 diset tpcc maria_user root
11 diset tpcc maria_pass 1234
12 diset tpcc maria_dbase tpcc
13
14 # Default for WH and VU
15 diset tpcc maria_count_ware 50
16 diset tpcc maria_num_vu 10
17
18 # Driver script options
19 diset tpcc maria_timeprofile true
20 diset tpcc maria_async_scale false
21 diset tpcc maria_driver timed
22
23 # Ensure test is limited by time
24 diset tpcc maria_total_iterations 1000000
25
26 # Timed duration
27 diset tpcc maria_rampup 2
28 diset tpcc maria_duration 8
29 # Distribute load
30 diset tpcc maria_allwarehouse false
31
32 # Transactions options
33 tcset refreshrate 10
34 tcset logtotemp 1
35 tcset unique 1
36 tcset timestamps 1
37
38 # Run
39 foreach z {2 4 8 12} {
40     set w [expr {$z * 5}]
41     diset tpcc maria_count_ware $w
42     puts "Building Schema for $z TEST"
43     # Delete possible previous data
44     deleteschema
45     vudestroy
46     # Build and load
47     buildschema
48     vudestroy
49     loadscript
50     # Vuser options
51     vuset delay 500
52     vuset repeat 500
53     vuset iterations 1
54     vuset showoutput 0
55     vuset logtotemp 1
56     vuset unique 1
57     vuset nobuff 0
58     vuset timestamps 1
59     puts "Starting $z VU TEST"
60     tcstart
61     vuset vu $z
62     vucreate
63     vurun
64     puts "Waiting 1 for cleanup and collection"
65     after 60000
66     puts "Destroying VU"
67     vudestroy
68     tcstop
69 }
```

A.3 PostgreSQL TCL Script

Listing 3: TCL script for PostgreSQL

```
1  #!/usr/bin/tclsh
2  # Set Database & Benchmark
3  dbset db pg
4  dbset bm TPC-C
5
6  # DB configs
7  diset connection pg_host 172.22.0.2
8  diset connection pg_port 5432
9  diset connection pg_sslmode prefer
10 diset tpcc pg_superuser postgres
11 diset tpcc pg_superuserpass 1234
12 diset tpcc pg_defaultdb postgres
13 diset tpcc pg_user tpcc
14 diset tpcc pg_pass tpcc
15 diset tpcc pg_dbase tpcc
16 diset tpcc pg_tspspace pg_default
17
18 # Default for WH and VU
19 diset tpcc pg_count_ware 50
20 diset tpcc pg_num_vu 10
21
22 # Driver script options
23 diset tpcc pg_timeprofile true
24 diset tpcc pg_async_scale false
25 diset tpcc pg_driver timed
26 # Ensure test is limited by time
27 diset tpcc pg_total_iterations 1000000
28 # Timed duration
29 diset tpcc pg_rampup 2
30 diset tpcc pg_duration 8
31 # Distribute load
32 diset tpcc pg_allwarehouse false
33 # Transactions options
34 tcset refreshrate 10
35 tcset logtotemp 1
36 tcset unique 1
37 tcset timestamps 1
38 # Run
39 foreach z {2 4 8 12} {
40     set w [expr {$z * 5}]
41     diset tpcc pg_count_ware $w
42     puts "Building Schema for $z TEST"
43     # Delete possible previous data
44     deleteschema
45     vudestroy
46     # Build and load
47     buildschema
48     vudestroy
49     loadscript
50     # Vuser options
51     vuset delay 500
52     vuset repeat 500
53     vuset iterations 1
54     vuset showoutput 0
55     vuset logtotemp 1
56     vuset unique 1
57     vuset nobuff 0
58     vuset timestamps 1
59     puts "Starting $z VU TEST"
60     tcstart
61     vuset vu $z
62     vucreate
63     vurun
64     puts "Waiting 1 for cleanup and collection"
65     after 60000
66     puts "Destroying VU"
67     vudestroy
68     tcstop
69 }
```

B Docker

B.1 Docker Compose File

Listing 4: Docker Compose file for the databases and HammerDB

```
1 version: "3.9"
2
3 services:
4   mysql:
5     build:
6       context: .
7       dockerfile: Dockerfiles/MySQL/Dockerfile
8     restart: always
9     container_name: MySQL_SBD
10    environment:
11      MYSQL_ROOT_PASSWORD: 1234
12    volumes:
13      - mysql-volume:/var/lib/mysql
14    networks:
15      sbd_network:
16        ipv4_address: 172.22.0.3
17
18   postgres:
19     image: postgres
20     container_name: Postgres_SBD
21     command: postgres -c config_file=/etc/postgresql/postgresql.conf
22     environment:
23       POSTGRES_PASSWORD: "1234"
24     volumes:
25       - postgres-volume:/var/lib/postgresql/data
26       - ./Configs/Hard_Test/postgresql.conf:/etc/postgresql/postgresql.conf:ro
27     networks:
28       sbd_network:
29         ipv4_address: 172.22.0.2
30
31   mariadb:
32     build:
33       context: .
34       dockerfile: Dockerfiles/MariaDB/Dockerfile
35     restart: always
36     container_name: MariaDB_SBD
37     environment:
38       MYSQL_ROOT_PASSWORD: 1234
39       MYSQL_DATABASE: sbdDatabase
40     volumes:
41       - mariadb-volume:/var/lib/mysql:rw
42     networks:
43       sbd_network:
44         ipv4_address: 172.22.0.4
45
46   hammer-gui:
47     image: tpcorg/hammerdb:latest-cloudtk
48     container_name: HammerDB_SBD
49     restart: always
50     ports:
51       - "8081:8081"
52       - "8082:8082"
53       - "8080:8080"
54     depends_on:
55       - mysql
56       - postgres
57       - mariadb
58     volumes:
59       - ./Scripts:/home/HammerDB-4.10/scripts/tcl-scripts
60       - ./stats-logs:/exp-logs
61     networks:
62       sbd_network:
63         ipv4_address: 172.22.0.5
64
65 volumes:
66   mysql-volume:
```

```
67     driver: local
68 postgres-volume:
69     driver: local
70 mariadb-volume:
71     driver: local
72
73 networks:
74     sbd_network:
75         driver: bridge
76         ipam:
77             config:
78                 - subnet: 172.22.0.0/16
```

B.2 Dockerfile for MySQL

Listing 5: Dockerfile for MySQL

```
1 FROM mysql:8.4
2
3 RUN mkdir -p /var/log/mysql && \
4     chown -R mysql:mysql /var/log/mysql
5
6 COPY ./Configs/Hard_Test/mysql.cnf /etc/mysql/conf.d/mysql.cnf
7
8 RUN chmod 644 /etc/mysql/conf.d/mysql.cnf
```

B.3 Dockerfile for MariaDB

Listing 6: Dockerfile for MariaDB

```
1 FROM mariadb:11.4
2
3 COPY ./Configs/Hard_Test/my.cnf /etc/mysql/my.cnf
4
5 RUN chmod 644 /etc/mysql/my.cnf
```

C Configs

C.1 MySQL Configuration File

Listing 7: MySQL configuration file

```
1 [mysqld]
2
3 #-----
4 # BASIC SETTINGS
5 #-----
6 port = 3306
7 bind-address = 0.0.0.0
8 max_connections = 1000
9 max_connect_errors = 1000000
10 wait_timeout = 28800
11 interactive_timeout = 28800
12 connect_timeout = 10
13 back_log = 1500
14
15 #-----
16 # CHARACTER SET & COLLATION
17 #-----
18 character-set-server = utf8mb4
19 collation-server = utf8mb4_unicode_ci
20
21 #-----
22 # STORAGE AND INNODB ENGINE
23 #-----
24 innodb_buffer_pool_size = 4G
25 innodb_buffer_pool_instances = 4
26 innodb_log_file_size = 512M
27 innodb_log_buffer_size = 64M
28 innodb_file_per_table = 1
29 innodb_flush_log_at_trx_commit = 2
30 innodb_flush_method = O_DIRECT
31 innodb_io_capacity = 2000
32 innodb_io_capacity_max = 4000
33 innodb_read_io_threads = 8
34 innodb_write_io_threads = 8
35 innodb_purge_threads = 4
36 innodb_doublewrite = 1
37 innodb_autoinc_lock_mode = 2
38 innodb_stats_persistent = 1
39 innodb_lru_scan_depth = 2048
40 innodb_adaptive_flushing = 1
41 innodb_adaptive_hash_index = 0
42 innodb_change_buffering = none
43
44 #-----
45 # TEMPORARY TABLE & BUFFERS
46 #-----
47 tmp_table_size = 256M
48 max_heap_table_size = 256M
49 sort_buffer_size = 1M
50 join_buffer_size = 1M
51 read_buffer_size = 512K
52 read_rnd_buffer_size = 2M
53
54 #-----
55 # LOGGING
56 #-----
57 slow_query_log = 1
58 long_query_time = 2
59 slow_query_log_file = /var/lib/mysql/mysql-slow.log
60 general_log = 0
61
62 #-----
63 # BINARY LOGGING
64 #-----
65 skip-log-bin
66 sync_binlog = 0
```

```

67 | #-----
68 | # SECURITY & COMPATIBILITY
69 | #-----
70 | local_infile = 0
71 | sql_mode = STRICT_TRANS_TABLES,ERROR_FOR_DIVISION_BY_ZERO,NO_ENGINE_SUBSTITUTION
72 | transaction_isolation = REPEATABLE-READ
73 |
74 | #-----
75 | # PERFORMANCE
76 | #-----
77 | performance_schema = OFF
78 | max_prepared_stmt_count = 12800
79 | table_open_cache = 2000
80 | table_open_cache_instances = 4
81 | open_files_limit = 65535
82 | thread_cache_size = 50
83 | thread_stack = 256K
84 |
85 | #-----
86 | # MONITORING
87 | #-----
88 | innodb_monitor_enable = '%'
89 |
90 | #-----
91 | # PLUGIN & AUTHENTICATION
92 | #-----
93 | require_secure_transport = OFF
94 | caching_sha2_password_auto_generate_rsa_keys = ON
95 |

```

C.2 MariaDB Configuration File

Listing 8: MariaDB configuration file

```

1 | [mysqld]
2 |
3 | #-----
4 | # BASIC SETTINGS
5 | #-----
6 | port = 3306
7 | bind-address = 0.0.0.0
8 | max_connections = 1000
9 | max_connect_errors = 1000000
10 | wait_timeout = 28800
11 | interactive_timeout = 28800
12 | connect_timeout = 10
13 | back_log = 1500
14 |
15 | #-----
16 | # CHARACTER SET & COLLATION
17 | #-----
18 | character-set-server = utf8mb4
19 | collation-server = utf8mb4_unicode_ci
20 |
21 | #-----
22 | # STORAGE AND INNODB ENGINE
23 | #-----
24 | innodb_buffer_pool_size = 4G
25 | innodb_buffer_pool_instances = 4
26 | innodb_log_file_size = 512M
27 | innodb_log_buffer_size = 64M
28 | innodb_file_per_table = 1
29 | innodb_flush_log_at_trx_commit = 2
30 | innodb_flush_method = O_DIRECT
31 | innodb_io_capacity = 2000
32 | innodb_io_capacity_max = 4000
33 | innodb_read_io_threads = 8
34 | innodb_write_io_threads = 8
35 | innodb_purge_threads = 4
36 | innodb_doublewrite = 1
37 | innodb_autoinc_lock_mode = 2

```

```

38 innodb_stats_persistent = 1
39 innodb_lru_scan_depth = 2048
40 innodb_adaptive_flushing = 1
41 innodb_adaptive_hash_index = 0
42 innodb_change_buffering = none
43
44 #-----
45 # TEMPORARY TABLE & BUFFERS
46 #-----
47 tmp_table_size = 256M
48 max_heap_table_size = 256M
49 sort_buffer_size = 1M
50 join_buffer_size = 1M
51 read_buffer_size = 512K
52 read_rnd_buffer_size = 2M
53
54 #-----
55 # LOGGING
56 #-----
57 slow_query_log = 1
58 long_query_time = 2
59 slow_query_log_file = /var/lib/mysql/mysql-slow.log
60 general_log = 0
61
62 #-----
63 # BYNARY LOGGING
64 #-----
65 skip-log-bin
66 sync_binlog = 0
67
68 #-----
69 # SECURITY & COMPATIBILITY
70 #-----
71 local_infile = 0
72 sql_mode = STRICT_TRANS_TABLES,ERROR_FOR_DIVISION_BY_ZERO,NO_ENGINE_SUBSTITUTION
73 transaction_isolation = REPEATABLE-READ
74
75 #-----
76 # PERFORMANCE
77 #-----
78 performance_schema = OFF
79 max_prepared_stmt_count = 12800
80 table_open_cache = 2000
81 table_open_cache_instances = 4
82 open_files_limit = 65535
83 thread_cache_size = 50
84 thread_stack = 256K
85
86 #-----
87 # MONITORING
88 #-----
89 innodb_monitor_enable = '%'

```


C.3 PostgreSQL Configuration File

Listing 9: PostgreSQL configuration file

```
1 #-----
2 # FILE LOCATIONS
3 #-----
4 data_directory = '/var/lib/postgresql/data'
5 hba_file = '/var/lib/postgresql/data/pg_hba.conf'
6 ident_file = '/var/lib/postgresql/data/pg_ident.conf'
7
8 #-----
9 # CONNECTIONS AND AUTHENTICATION
10 #-----
11 listen_addresses = '*'
12 port = 5432
13 max_connections = 100
14
15 #-----
16 # RESOURCE USAGE
17 #-----
18 shared_buffers = 512MB
19 work_mem = 64MB
20 maintenance_work_mem = 256MB
21 effective_cache_size = 2GB
22
23 #-----
24 # WRITE-AHEAD LOG
25 #-----
26 wal_level = replica
27 synchronous_commit = off
28 checkpoint_timeout = 15min
29 checkpoint_completion_target = 0.9
30 max_wal_size = 2GB
31 min_wal_size = 512MB
32
33 #-----
34 # LOGGING (optional, but useful for debugging)
35 #-----
36 logging_collector = on
37 log_directory = 'log'
38 log_filename = 'postgresql.log'
39 log_statement = 'none'
40 log_min_duration_statement = 1000
```

D Runner Scripts

D.1 Batch Script

Listing 10: Batch script to run the TCL scripts

```
1 @echo off
2 setlocal
3
4 :: Set container names
5 set HAMMER_CONTAINER=HammerDB_SBD
6
7 echo Starting HammerDB benchmark runner...
8
9 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
10 :: Setup containers
11 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
12
13 :: Compose down to ensure a clean start
14 echo Stopping and removing existing containers, images and volumes...
15 docker compose down --rmi all -v
16
17 timeout /t 5
18
19 :: Compose up to start the containers
20 echo Starting containers...
21 docker compose up -d
22
23 timeout /t 10
24
25 echo Containers are up and running!
26
27 echo Starting HammerDB benchmark...
28
29 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
30 :: POSTGRESQL
31 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
32
33 echo Starting up PostgreSQL benchmark...
34
35 :: Start the stats logging in a new window using the temp file
36 start "PCStats" cmd /c PCStats.bat "postgres"
37
38 timeout /t 5 >nul
39
40 docker exec -i %HAMMER_CONTAINER% /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/largeTestPG.tcl
41
42 for /f "tokens=2 delims=," %i in ('tasklist /v /fo csv ^| findstr /i /c:"PCStats"') do (
43     taskkill /PID %i /F >nul 2>&1
44 )
45
46 echo Benchmark and monitoring complete for PostgreSQL.
47
48 timeout /t
49
50 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
51 :: MYSQL
52 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
53
54 echo Starting up MySQL benchmark...
55
56 :: Start the stats logging in a new window using the temp file
57 start "PCStats" cmd /c PCStats.bat "mysql"
58
59 timeout /t 5 >nul
60
61 docker exec -i %HAMMER_CONTAINER% /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/largeTestMySQL.
62     ↪ tcl
63
64 for /f "tokens=2 delims=," %i in ('tasklist /v /fo csv ^| findstr /i /c:"PCStats"') do (
65     taskkill /PID %i /F >nul 2>&1
66 )
```

```

66
67 echo Benchmark and monitoring complete for MySQL.
68
69 timeout /t 30
70
71 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
72 :: MARIADB
73 ::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::::
74
75 echo Starting up MariaDB benchmark...
76
77 :: Start the stats logging in a new window using the temp file
78 start "PCStats" cmd /c PCStats.bat "mariadb"
79
80 timeout /t 5 >nul
81
82 docker exec -i %HAMMER_CONTAINER% /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/
    ↪ largeTestMariaDB.tcl
83
84 for /f "tokens=2 delims=" %%i in ('tasklist /v /fo csv ^| findstr /i /c:"PCStats"') do (
85     taskkill /PID %%~i /F >nul 2>&1
86 )
87
88 echo Benchmark and monitoring complete for MariaDB.
89
90 echo Benchmark and monitoring complete.
91 endlocal

```

D.1.1 PCStats Script

Listing 11: PCStats script to monitor the system in Windows

```

1 @echo off
2 setlocal enabledelayedexpansion
3
4 set DB=%1
5
6 :: Format date and time safely for filename
7 for /f "tokens=1-3 delims=/" %%a in ("%date%") do (
8     set mm=%%a
9     set dd=%%b
10    set yyyy=%%c
11 )
12 for /f "tokens=1-3 delims=:" %%a in ("%time%") do (
13     set hh=%%a
14     set min=%%b
15     set ss=%%c
16 )
17 set "logfile=sys_usage_%mm%_%dd%_%hh%_%min%_%DB%.csv"
18
19 :: Set interval (in seconds)
20 set interval=10
21
22 :: Write CSV header
23 echo Timestamp,CPU_Usage_Percent,RAM_Used_MB > "%logfile%"
24
25 :loop
26 :: Get timestamp in ISO format using PowerShell
27 for /f %%x in ('powershell -command "Get-Date -Format yyyy-MM-dd_HH:mm:ss"') do set timestamp=%%x
28
29 echo Current timestamp: %timestamp%
30
31 :: Get CPU usage
32 for /f "skip=1" %%x in ('wmic cpu get loadpercentage') do (
33     if not "%%x"==" " (
34         set cpu=%%x
35         goto :gotCPU
36     )
37 )
38 :gotCPU
39

```

```

40 :: Get RAM usage (in MB)
41 for /f "skip=1 tokens=2,3 delims=" %%a in ('wmic OS get FreePhysicalMemory^,TotalVisibleMemorySize /
    ↪ format:csv') do (
42     set "free=%%a"
43     set "total=%%b"
44 )
45 set /a usedMB=(%total% - %free%) / 1024
46
47 :: Write to CSV
48 echo %timestamp%,%cpu%,%usedMB% >> "%logfile%"
49
50 :: Wait and repeat
51 timeout /t %interval% >nul
52 goto loop

```

D.2 Shell Script

Listing 12: Shell script to run the TCL scripts

```

1  #!/bin/bash
2
3  # Set your container names
4  HAMMER_CONTAINER=HammerDB_SBD
5
6  echo "Starting HammerDB benchmark runner..."
7
8  #####
9  # Setup containers
10 #####
11
12 # Compose down to ensure a clean start
13 echo "Stopping and removing existing containers, images and volumes..."
14
15 docker compose down --rmi all -v
16
17 sleep 5
18
19 # Compose up to start the containers
20 echo "Starting containers..."
21
22 docker compose up -d
23
24 sleep 10
25
26 echo "Containers are up and running!"
27
28 echo "Starting HammerDB benchmark..."
29
30 #####
31 # POSTGRESQL
32 #####
33
34 echo "Starting up PostgreSQL benchmark..."
35
36 # Start the stats logging in the background
37 ./PCStats.sh "postgres" &
38 PCSTATS_PID=$!
39
40 # Wait 5 seconds for stats logger to initialize
41 sleep 5
42
43 # Run the benchmark via Docker
44 docker exec -i "$HAMMER_CONTAINER" /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/largeTestPG.
    ↪ tcl
45
46 # Kill the PCStats logger
47 kill "$PCSTATS_PID"
48
49 echo "Benchmark and monitoring complete for PostgreSQL."
50
51 # Wait 30 seconds before exiting

```

```

52 sleep 30
53
54 #####
55 # MYSQL
56 #####
57
58 echo "Starting up MySQL benchmark..."
59
60 # Start the stats logging in the background
61 ./PCStats.sh "mysql" &
62 PCSTATS_PID=$!
63
64 # Wait 5 seconds for stats logger to initialize
65 sleep 5
66
67 # Run the benchmark via Docker
68 docker exec -i "$HAMMER_CONTAINER" /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/largeTestMySQL
    ↪ .tcl
69
70 # Kill the PCStats logger
71 kill "$PCSTATS_PID"
72
73 echo "Benchmark and monitoring complete for MySQL."
74
75 # Wait 30 seconds before exiting
76 sleep 30
77
78 #####
79 # MARIADB
80 #####
81
82 echo "Starting up MariaDB benchmark..."
83
84 # Start the stats logging in the background
85 ./PCStats.sh "mariadb" &
86 PCSTATS_PID=$!
87
88 # Wait 5 seconds for stats logger to initialize
89 sleep 5
90
91 # Run the benchmark via Docker
92 docker exec -i "$HAMMER_CONTAINER" /home/HammerDB-4.10/hammerdbcli auto scripts/tcl-scripts/
    ↪ largeTestMariaDB.tcl
93
94 # Kill the PCStats logger
95 kill "$PCSTATS_PID"
96
97 echo "Benchmark and monitoring complete for MariaDB."
98
99 echo "Benchmark and monitoring complete."

```

D.2.1 PCStats Script

Listing 13: PCStats script to monitor the system in Linux

```

1 #!/bin/bash
2
3 DB="$1"
4 interval=10
5
6 # Create log file name based on current date and time
7 timestamp=$(date "+%m_%d_%H_%M")
8 logfile="sys_usage_${timestamp}_${DB}.csv"
9
10 # Write CSV header
11 echo "Timestamp,CPU_Usage_Percent,RAM_Used_MB" > "$logfile"
12
13 while true; do
14     # Get current timestamp
15     timestamp=$(date "+%Y-%m-%d_%H:%M:%S")
16     echo "Current timestamp: $timestamp"

```

```

17
18 # Get CPU usage percentage (average over 1 second)
19 cpu=$(top -bn2 | grep "Cpu(s)" | tail -n 1 | awk -F'id,' -v prefix="" '{ split($1, vs, ","); cpu=100 -
    ↪ vs[length(vs)]; printf "%.0f", cpu }')
20
21 # Get RAM usage in MB
22 mem_total=$(free -m | awk '/Mem:/ {print $2}')
23 mem_used=$(free -m | awk '/Mem:/ {print $3}')
24
25 # Write to CSV
26 echo "$timestamp,$cpu,$mem_used" >> "$logfile"
27
28 # Wait for interval
29 sleep "$interval"
30 done

```