



NOVA SCHOOL OF  
SCIENCE & TECHNOLOGY

NOVA UNIVERSITY OF LISBON

MSC IN COMPUTER SCIENCE

# **PERFORMANCE COMPARISON BETWEEN DBMSs UNDER TPROC-C WORKLOADS**

*José Costa (62637)*  
*Rodrigo Albuquerque (70294)*  
*Rodrigo Silva (70567)*

DATABASES SYSTEMS

JUNE 4, 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Overview of HammerDB</b>	<b>1</b>
2.1	Overview of TPROC-C . . . . .	1
2.2	TPROC-C vs TPROC-H . . . . .	1
<b>3</b>	<b>Problem &amp; DBMS Summary</b>	<b>1</b>
<b>4</b>	<b>Benchmark Description</b>	<b>1</b>
<b>5</b>	<b>Methodology</b>	<b>1</b>
5.1	Hardware and Software Setup . . . . .	2
5.2	Database Setup . . . . .	2
<b>6</b>	<b>Results</b>	<b>3</b>
<b>7</b>	<b>Discussion</b>	<b>3</b>
<b>8</b>	<b>Conclusions</b>	<b>3</b>

# 1 Introduction

This project aims to benchmark and compare the performance of different database systems using the TPC-C workload, a standard for evaluating OLTP (Online Transaction Processing) environments.

Automated scripts are used to run identical tests across multiple databases, each configured with similar settings to ensure a fair comparison.

Key metrics such as Transactions Per Minute (TPM) are measured under varying levels of concurrent users.

The results help identify the strengths and limitations of each database in handling transactional workloads.

In this test, the databases used were PostgreSQL, MySQL and MariaDB.

In total, we ran 54 tests:

- 4 tests scaling the number of virtual users (2 4 8 12) and warehouses (VU\*5) on all PCs and databases (48 tests in total);
- 1 test with the number of virtual users set to the same as the number of threads in that PC, warehouses set to VU\*5, allwarehouse = true on all databases on just one PC (3 test in total);
- 1 test with the number of virtual users set to the same as the number of threads in that PC, warehouses set to VU\*5, with default config on all databases on just one PC (3 test in total).

## 2 Overview of HammerDB

HammerDB is a free, open-source tool for benchmarking the performance of relational databases [1].

It supports popular databases like Oracle, SQL Server, PostgreSQL, MySQL, and more. HammerDB uses industry-standard workloads such as TPROC-C and TPROC-H to simulate real-world database activity.

It offers both a graphical interface and command-line options, making it suitable for developers, DBAs, and system administrators to test, compare, and tune database performance.

In some cases we used HammerDB in docker containers to run the tests, which allows for easy setup and isolation of the testing environment.

In another case, we used the Windows version of HammerDB to run the tests on a Windows machine.

### 2.1 Overview of TPROC-C

TPROC-C is a benchmark designed to evaluate the performance of database management systems (DBMS) using a transactional workload. It simulates a typical online transaction processing (OLTP) environment, focusing on operations like inserts, updates, and deletes across multiple tables.

### 2.2 TPROC-C vs TPROC-H

TPROC-H is a benchmark designed for data warehousing and analytical workloads, while TPROC-C is focused on transactional processing. TPROC-H emphasizes complex queries and large data sets, whereas TPROC-C simulates real-time transactions with a focus on insert, update, and delete operations.

## 3 Problem & DBMS Summary

## 4 Benchmark Description

## 5 Methodology

## 5.1 Hardware and Software Setup

PC	1	2	3	4
OS	Windows 11	Windows 11	Linux (Unraid)	MacOS Sequoia
CPU	Intel i7-13700H	AMD Ryzen 5 3600	Intel i3-10100F	Apple M1
Cores	14 (6P 8E)	6	4	8
Threads	20	12	8	8
RAM	16GB	16GB	32GB	16GB
Disk	SSD M.2 NVMe	SSD M.2 NVMe	SSD M.2 NVMe	SSD M.2 NVMe
Read	3500 MB/s	2500 MB/s	3500 MB/s	3400 MB/s
Write	2700 MB/s	2100 MB/s	3300 MB/s	2800 MB/s

Table 1: Hardware used in the benchmarks

In PCs **1**, **3**, and **4**, we used HammerDB and all the databases in docker containers to run the tests, which allows for easy setup and isolation of the testing environment.

We used docker compose for this setup, which allowed us to easily run the same tests on different machines with the same configuration.

In PC **2**, we used everything installed on the host machine.

## 5.2 Database Setup

For MySQL and MariaDB, this was the configuration used:

```
1 [mysqld]
2
3 # BASIC SETTINGS
4 port = 3308
5 bind-address = 0.0.0.0
6 max_connections = 1000
7 max_connect_errors = 1000000
8 wait_timeout = 28800
9 interactive_timeout = 28800
10 connect_timeout = 10
11 back_log = 1500
12
13 # CHARACTER SET & COLLATION
14 character-set-server = utf8mb4
15 collation-server = utf8mb4_unicode_ci
16
17 # STORAGE AND INNODB ENGINE
18 innodb_buffer_pool_size = 4G
19 innodb_buffer_pool_instances = 4
20 innodb_log_file_size = 512M
21 innodb_log_buffer_size = 64M
22 innodb_file_per_table = 1
23 innodb_flush_log_at_trx_commit = 2
24 innodb_flush_method = O_DIRECT
25 innodb_io_capacity = 2000
26 innodb_io_capacity_max = 4000
27 innodb_read_io_threads = 8
28 innodb_write_io_threads = 8
29 innodb_purge_threads = 4
30 innodb_doublewrite = 1
31 innodb_autoinc_lock_mode = 2
32 innodb_stats_persistent = 1
33 innodb_lru_scan_depth = 2048
34 innodb_adaptive_flushing = 1
35 innodb_adaptive_hash_index = 0
36 innodb_change_buffering = none
37
38 # TEMPORARY TABLE & BUFFERS
39 tmp_table_size = 256M
40 max_heap_table_size = 256M
41 sort_buffer_size = 1M
42 join_buffer_size = 1M
43 read_buffer_size = 512K
```

```

44 read_rnd_buffer_size = 2M
45
46 # LOGGING
47 slow_query_log = 1
48 long_query_time = 2
49 slow_query_log_file = /var/lib/mysql/mysql-slow.log
50 general_log = 0
51
52 # BINARY LOGGING
53 skip-log-bin
54 sync_binlog = 0
55
56 # SECURITY & COMPATIBILITY
57 local_infile = 0
58 sql_mode = STRICT_TRANS_TABLES,ERROR_FOR_DIVISION_BY_ZERO,NO_ENGINE_SUBSTITUTION
59 transaction_isolation = REPEATABLE-READ
60
61 # PERFORMANCE
62 performance_schema = OFF
63 max_prepared_stmt_count = 12800
64 table_open_cache = 2000
65 table_open_cache_instances = 4
66 open_files_limit = 65535
67 thread_cache_size = 50
68 thread_stack = 256K
69
70 # MONITORING
71 innodb_monitor_enable = '%'

```

For PostgreSQL, this was the configuration used:

```

1 # FILE LOCATIONS
2 data_directory = '/var/lib/postgresql/data' # Important: match Docker volume
3 hba_file = '/var/lib/postgresql/data/pg_hba.conf'
4 ident_file = '/var/lib/postgresql/data/pg_ident.conf'
5
6 # CONNECTIONS AND AUTHENTICATION
7 listen_addresses = '*' # Allow external connections (Docker host network)
8 port = 5432 # Default PostgreSQL port
9 max_connections = 100
10
11 # RESOURCE USAGE
12 shared_buffers = 512MB # Adjust depending on host memory (e.g., 25% of RAM)
13 work_mem = 64MB # Suitable for OLTP like TPC-C
14 maintenance_work_mem = 256MB
15 effective_cache_size = 2GB # Depends on total system RAM
16
17 # WRITE-AHEAD LOG
18 wal_level = replica
19 synchronous_commit = off # Can improve write performance (acceptable for benchmarks)
20 checkpoint_timeout = 15min
21 checkpoint_completion_target = 0.9
22 max_wal_size = 2GB
23 min_wal_size = 512MB
24
25 # LOGGING
26 logging_collector = on
27 log_directory = 'log'
28 log_filename = 'postgresql.log'
29 log_statement = 'none'
30 log_min_duration_statement = 1000 # Log queries slower than 1s

```

## 6 Results

## 7 Discussion

## 8 Conclusions

## Bibliography

- [1] Wikipedia contributors. *HammerDB — Wikipedia, The Free Encyclopedia*. [Online; accessed 25-May-2025]. 2025. URL: <https://en.wikipedia.org/w/index.php?title=HammerDB&oldid=1275860580>.